

Assignment 2.

This Assignment consist of 2 parts. First part related with KNN, second is K-Means Clusteing.

Part 1. KNN.

In this task you are going to predict text.csv with KNN algorithm. You will work with Breast Cancer Wisconsin (Diagnostic) Data Set. More info [here](#).

The data set is already spited to the train and test files. On the moodle you can find 4 files(train.csv, test.csv, train_answers.csv, test_answers.csv). Please open the files and start to explore. train.csv and test.csv contains 31 columns, here is the description of the 10 columns:

- a) radius (mean of distances from center to points on the perimeter)
- b) texture (standard deviation of gray-scale values)
- c) perimeter
- d) area
- e) smoothness (local variation in radius lengths)
- f) compactness ($\text{perimeter}^2 / \text{area} - 1.0$)
- g) concavity (severity of concave portions of the contour)
- h) concave points (number of concave portions of the contour)
- i) symmetry
- j) fractal dimension ("coastline approximation" - 1)
- etc.....

train_answers.csv and test_answers.csv they labels for the train and test set. There is two columns id and Diagnosis (m=malignant, b=benign)

Dataset is Standardized no need to scale them. Do not touch the test_answers.csv at the beginning . You will use it at the final step, when you start to calculate the accuracy score.

Step to do:

Please open the files

```
import pandas as pd

train = pd.read_csv("train.csv", index_col='Unnamed: 0')
y = pd.read_csv("train_answers.csv", index_col='Unnamed: 0')
test = pd.read_csv("test.csv", index_col='Unnamed: 0')
answers = pd.read_csv("real_answers_for_test.csv", index_col='Unnamed: 0')
```

Look at their sizes with data.shape. How many features we have? What is the sizes of train, test, y and answers?

Start to predicting test with help of train and y. **NOTE: you will not implement KNN algorithm from packages. For the calculating the Euclidean distance you can use `numpy.linalg.norm(a-b)`.**

Example of predicting first row of test with one neighbor:

- I. Calculate the distance to all train set.

```
# record the distance
distLst = []
for i in range(len(train)):
    #calcualte the distance
    distance = np.linalg.norm(test.iloc[0]-train.iloc[i])
    # record the distance
    distLst.append(distance)
```


II. Find the Minimum distance and predict to according it's label.

```
print("Minimum distance: ", min(distLst))
print("The index of the record with minimum dist: ", distLst.index(min(distLst)))
print("The label at that index: ", y.iloc[13])
```

```
Minimum distance: 3.06778195989615
The index of the record with minimum dist: 13
The label at that index: Diagnosis B
```

In that example above we predicted that first example of test is B (benign).

Your task is to predict for all other records of test with 1, 3, 5, 6, 7, 9 nearest neighbors. Calculate the accuracy score (without using package) for each neighbors result with test_answers.csv. and plot the graph. Accuracy VS number of the neighbors.



This the the result that you should get finally.

Part 2. K-means.

In this part you will create two functions. We know that we have two important steps in iteration:

1. Finding the closest centroid (assign it)
2. Compute the mean. (recomputing at each iteration)

Algorithm 1 k -means algorithm

- 1: Specify the number k of clusters to assign.
 - 2: Randomly initialize k centroids.
 - 3: **repeat**
 - 4: **expectation:** Assign each point to its closest centroid.
 - 5: **maximization:** Compute the new centroid (mean) of each cluster.
 - 6: **until** The centroid positions do not change.
-

```
for i in range(iterations):

    idx = findClosestCentroids(X, centroids)

    |
    centroids = computeMeans(X, idx, K)
```

Your task is to create function findClosestCentroids.

```
def findClosestCentroids(X, centroids):  
    .. ..
```

X is data matrix, centroids it is location of all centroids. Output of should be array that holds the index (a value in $\{1, \dots, K\}$ K is total number of centroids) of the closest centroid to every training example.

This is the result of function :

```
# Load an example dataset that we will be using  
data = loadmat('ex7data2.mat')  
X = data['X']  
initial_centroids = np.array([[3, 3], [6, 2], [8, 5]])  
  
# Find the closest centroids for the examples using the initial_centroids  
idx = findClosestCentroids(X, initial_centroids)  
  
print('Closest centroids for the first 3 examples:')  
print(idx[:3])  
print('(the closest centroids should be 0, 2, 1 respectively)')
```

Closest centroids for the first 3 examples:

```
[[0]  
 [2]  
 [1]]
```

(the closest centroids should be 0, 2, 1 respectively)

Next task is is to create function computeCentroids.

```
def computeCentroids(X, idx, K):
```

idx is the A vector (size m) of centroid assignments (i.e. each entry in range $[0 \dots K-1]$) example. K is the number of clusters. and you should return centroids A matrix of size (K, n) , n is the number of features.

This is the result of function :

```
centroids = computeCentroids(X, idx, K)  
  
print('Centroids computed after initial finding of closest centroids:')  
print(centroids)  
print('\nThe centroids should be')  
print('    [ 2.428301 3.157924 ]')  
print('    [ 5.813503 2.633656 ]')  
print('    [ 7.119387 3.616684 ]')
```

Centroids computed after initial finding of closest centroids:

```
[[2.42830111 3.15792418]  
 [5.81350331 2.63365645]  
 [7.11938687 3.6166844 ]]
```

The centroids should be

```
[ 2.428301 3.157924 ]  
[ 5.813503 2.633656 ]  
[ 7.119387 3.616684 ]
```

If you correctly completed that two function above,
please run the following code

```
from scipy.io import loadmat
import os
import numpy as np

data = loadmat(os.path.join('Data', 'ex7data2.mat'))
X = data['X']
max_iters = 10
initial_centroids = np.array([[3, 3], [6, 2], [8, 5]])

for i in range(max_iters):
    idx = findClosestCentroids(X, centroids)
    centroids = computeCentroids(X, idx, K)
```

And plot the results.

